



Learning Objectives

- Understand the Cartesian product.
- Understand the concepts of joins.
- Understanding the different types of joins.
- Understanding of combining data across multiple joins.

LO1: Understanding Cartesian product.

Cartesian Product

The Cartesian Product in SQL is a type of join that returns the combination of every row from one table with every row from another table. It is also called a **Cross Join** and results in an $m \times n$ combination of rows, where m is the number of rows in the first table and n is the number of rows in the second table.

Example: Let us consider the following tables:

customers

CustomerId	ContactName	City
1	Ali	Lahore
2	Ahmed	Faisalabad
3	Aslam	Karachi

orders

OrderId	CustomerId	ShipAddress	ShipPostalCode
101	1	Lahore	20022
102	2	Karachi	30011
103	2	Lahore	15022

Query:

```
SELECT *  
FROM customers, orders;
```

Result:

CustomerId	ContactName	City	CustomerId	OrderId	ShipAddress	ShipPostalCode
1	Ali	Lahore	1	101	Lahore	20022
1	Ali	Lahore	2	102	Karachi	30011
1	Ali	Lahore	2	103	Lahore	15022
2	Ahmed	Faisalabad	1	101	Lahore	20022
2	Ahmed	Faisalabad	2	102	Karachi	30011
2	Ahmed	Faisalabad	2	103	Lahore	15022
3	Aslam	Karachi	1	101	Lahore	20022
3	Aslam	Karachi	2	102	Karachi	30011
3	Aslam	Karachi	2	103	Lahore	15022

Cartesian Product with a Condition

To avoid unnecessary combinations, we can apply a condition to filter only the relevant rows. This is done using the **WHERE** clause or a **JOIN** condition.

Example: Filtering Based on Customer ID

Query:

```
SELECT *  
FROM customers, orders  
WHERE customers.CustomerId = orders.CustomerId;
```

Result:

CustomerId	ContactName	city	OrderId	shipAddress	ShipPostalCode
1	Ali	Lahore	101	Lahore	20022
2	Ahmed	Faisalabad	102	Karachi	30011
2	Ahmed	Faisalabad	103	Lahore	15022

LO2: Understanding the Concept of Joins.

Joins in SQL Server

A JOIN clause is used to combine rows from two or more tables, based on a related column between them.

Relationship Between Cartesian Product and Joins

In SQL, **joins** are used to combine rows from two or more tables based on a specified condition. The **Cartesian product** (also known as the **cross join**) is the foundation of all types of joins. Joins work by **first**

creating a Cartesian product and then filtering the rows based on a specified condition.

Join = Cartesian Product + Condition

Joins Syntax

```
SELECT column_name
FROM table1
INNER JOIN/ CROSS JOIN/ NATURAL JOIN/ LEFT JOIN/ RIGHT JOIN
table2
ON table1.column_name = table2.column_name;
```

LO3: Understanding the Different types of Joins

SQL server provides the following joins:

- Natural Join
- Inner Join
- Left Outer Join
- Right Outer Join
- Full Outer Join

Natural Join

A **Natural Join** is used when you want to combine two tables based on all columns with the same name and compatible data types, without explicitly specifying the join condition.

Example: Let us consider the below tables first

customers

CustomerId	ContactName	City
1	Ali	Lahore
2	Ahmed	Faisalabad
3	Aslam	Karachi

orders

OrderId	CustomerId	ShipAddress	ShipPostalCode
101	1	Lahore	200
102	2	Karachi	300
103	2	Lahore	150

Requirement: We have to report all customers who ordered.

Query:

```
SELECT *  
FROM customers  
NATURAL JOIN orders;
```

Result:

customerid	contactname	city	orderid	shipaddress	shippostalcode
1	Ali	Lahore	101	Lahore	200
2	Ahmed	Faisalabad	102	Karachi	150
2	Ahmed	Faisalabad	103	Lahore	150

Inner Join

An **Inner Join** is used to combine two tables based on a specified condition, usually matching values in a common column. Only rows that satisfy the condition (i.e., have matching values in both tables) are included in the result.

Example: Let us consider the tables below first

customers

CustomerId	ContactName	City
1	Ali	Lahore
2	Ahmed	Faisalabad
3	Aslam	Karachi

orders

orderid	customerid	shippingaddress	Shippostalcode
101	1	Lahore	200
102	2	Karachi	300
103	2	Lahore	150
104	4	Islamabad	400

Requirement: Generate a report listing only those customers who have made at least one purchase, along with their order details.

Query:

```
SELECT*  
FROM customers  
INNER JOIN orders  
ON customers.customerid = orders.customerid;
```

Result:

customer_id	contactname	city	orderid	shipaddress	shippostalcode
1	Ali	Lahore	101	Lahore	200
2	Ahmed	Faisalabad	102	Karachi	150
2	Ahmed	Faisalabad	103	Lahore	150

Left Outer Join

Left Outer join gives the matching rows and the rows which are in left table but not in right table.

Example: Let us consider the below tables first

customers

CustomerId	ContactName	City
1	Ali	Lahore
2	Ahmed	Faisalabad
3	Aslam	Karachi

orders

orderid	customerid	shipaddress	shippostalcode
101	1	Lahore	20022
102	2	Karachi	30011
103	2	Lahore	15022

Requirement: We have to report order details of customers and who does not give any orders show null orders also.

Query:

```
SELECT *
FROM customers C LEFT JOIN orders O
ON C.customerid = O.customerid
```

Result:

customerid	contactname	city	orderid	shipaddress	shippostalcode
1	Ali	Lahore	101	Lahore	20022
2	Ahmed	Faisalabad	102	Karachi	30011
2	Ahmed	Faisalabad	103	Lahore	15022
3	Aslam	Karachi	NULL	NULL	NULL

Right Outer Join

Right Outer join gives the matching rows and the rows which are in right table but not in left table.

Example: Let us consider the below tables first

customers

CustomerId	ContactName	City
1	Ali	Lahore
2	Ahmed	Faisalabad
3	Aslam	Karachi

orders

orderid	customerid	shipaddress	shippostalcode
101	1	Lahore	20022
102	2	Karachi	30011
103	2	Lahore	15022
104	4	Islamabad	40012

Requirement: We need to report all the details of orders, including information about the customers who placed those orders. Additionally, if an order exists without a corresponding customer in the customers table (due to missing or incorrect data), include these orders as well, showing NULL for the customer-related columns.

Query:

```
SELECT *
FROM customers C RIGHT JOIN orders O
```

ON C.customerid = O.customerid

Result:

customerId	Contactname	city	orderid	Shipaddress	shippostalcode
1	Ali	Lahore	101	Lahore	20022
2	Ahmed	Faisalabad	102	Karachi	30011
2	Ahmed	Faisalabad	103	Lahore	15022
Null	Null	Null	104	Islamabad	40012

Full Outer Join

Full Outer Join combines the results of both Left Outer Join and Right Outer Join:

- It includes **all rows from both tables**, even if there are no matches between them.
- If a match is found between the two tables, it combines the data into a single row

Example: Let us consider the below tables first

customers

CustomerId	ContactName	City
1	Ali	Lahore
2	Ahmed	Faisalabad
3	Aslam	Karachi

orders

orderid	customerid	shipaddress	shippostalcode
101	1	Lahore	20022
102	2	Karachi	30011
103	2	Lahore	15022
104	4	Islamabad	40012

Requirement: Give all details of customers and orders

Query:

```
SELECT *  
FROM customers C
```

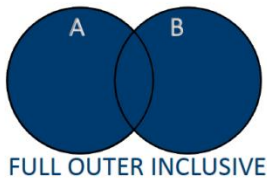
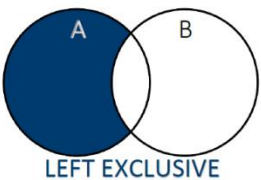
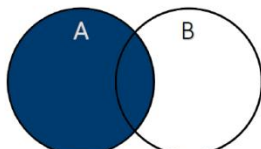
```

LEFT JOIN orders O ON C.customerId = O.customerId
UNION
SELECT *
FROM customers C
RIGHT JOIN orders O ON C.customerId = O.customerId;

```

Result:

customerId	Contactname	city	orderId	shipaddresses	ShipPostalCode
1	Ali	Lahore	101	Lahore	20022
2	Ahmed	Faisalabad	102	Karachi	30011
2	Ahmed	Faisalabad	103	Lahore	15022
3	Aslam	Karachi	Null	Null	Null
Null	Null	Null	104	Islamabad	40012



SQL JOINS	
LEFT INCLUSIVE SELECT [Select List] FROM TableA A LEFT OUTER JOIN TableB B ON A.Key= B.Key	RIGHT INCLUSIVE SELECT [Select List] FROM TableA A RIGHT OUTER JOIN TableB B ON A.Key= B.Key
LEFT EXCLUSIVE SELECT [Select List] FROM TableA A LEFT OUTER JOIN TableB B ON A.Key= B.Key WHERE B.Key IS NULL	RIGHT EXCLUSIVE SELECT [Select List] FROM TableA A LEFT OUTER JOIN TableB B ON A.Key= B.Key WHERE A.Key IS NULL
FULL OUTER INCLUSIVE SELECT [Select List] FROM TableA A FULL OUTER JOIN TableB B ON A.Key = B.Key	FULL OUTER EXCLUSIVE SELECT [Select List] FROM TableA A FULL OUTER JOIN TableB B ON A.Key = B.Key WHERE A.Key IS NULL OR B.Key IS NULL
INNER JOIN SELECT [Select List] FROM TableA A INNER JOIN TableB B ON A.Key = B.Key	

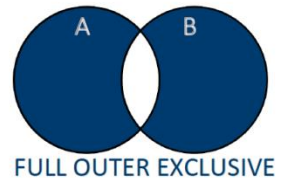
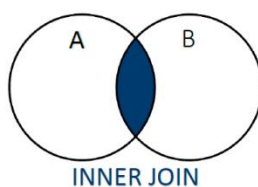
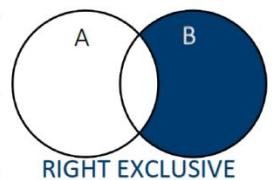
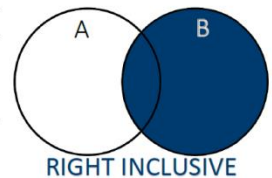


Figure 1 : Visual Representation of All Joins

LO4: Understanding of Combining Data across Multiple Joins

In SQL, complex queries often require combining data from multiple tables using different types of **JOINS**. By effectively using multiple joins, we can extract meaningful relationships and insights from relational databases.

Concept

When multiple tables are involved, we can use **multiple joins** to connect them based on related columns.

Example: Multiple Joins in Action

Let's consider a scenario where we need to fetch order details along with customer and product information.

customers

CustomerId	ContactName	city
1	Ali	Lahore
2	Aslam	Karachi

orders

orderId	customerId	product_id	amount
101	1	10	200
102	2	11	150

products

product_id	product_name	Price
10	Laptop	500
11	Phone	300

Requirement:

Retrieve customer names, cities, order amounts, and product names for all placed orders.

Query:

```
SELECT *  
FROM customers  
INNER JOIN orders ON customers.customerid = orders.customerid
```

```
INNER JOIN products ON orders.product_id =  
products.product_id;
```

Result:

name	city	amount	product_name
Ali	Lahore	200	Laptop
Aslam	Karachi	150	Phone

Tasks:

Perform all JOIN queries on any table using Northwind Schema. You cannot use group by clause in this manual.

1. Write a query that returns all possible combinations of **employees** and **shippers**.

Output: EmployeeID (EmployeeID), EmployeeName (EmployeeName), ShipperID (ShipperID), ShipperName (ShipperName)

2. Generate all possible **customer-order-shipper** combinations but filter only those where the **customer's first three characters matches the shipper's name first three characters**.

Output: CustomerID (CustomerID), CustomerName (CustomerName), ShipperID (ShipperID), ShipperName (ShipperName), First three characters

3. Create a Cartesian product of employees with themselves to identify all possible **employee pairs**. Ensure that each pair appears only once.

Output: Employee1ID (Employee1ID), Employee1Name (Employee1Name), Employee2ID (Employee2ID), Employee2Name (Employee2Name)

4. Display order details showing OrderID, CustomerName, ProductName, Quantity, Employee's Full Name, and Shipper's Name.

Output: OrderID (OrderID), CustomerName (CustomerName), ProductName (ProductName), Quantity (Quantity), EmployeeName (EmployeeName), ShipperName (ShipperName)

5. Retrieve a list of **all products** that have never been ordered

Output: ProductID (ProductID), ProductName (ProductName), CategoryName (CategoryName), SupplierName (SupplierName)

6. Retrieve a list of **employees** who have **not processed any orders in the last 6 months of an year**.
Output: EmployeeID (EmployeeID), EmployeeName (EmployeeName)
7. Retrieve a list of **all orders** including those that do not have an assigned **employee** or **customer**.
Output: OrderID (OrderID), CustomerID (CustomerID), CustomerName (CustomerName), EmployeeID (EmployeeID), EmployeeName (EmployeeName)
8. Retrieve all customers and their orders. If the customer has never placed an order, display "**No Order Placed**" in the OrderID column.
Output: CustomerID (CustomerID), CustomerName (CustomerName), OrderID (OrderID)
9. Find all orders placed **before 9 AM or after 5 PM**.
Output: OrderID (OrderID), CustomerID (CustomerID), CustomerName (CustomerName), OrderHour (OrderHour)
10. Identify **pairs of employees who have both processed orders for the same product**.
Output: ProductID (ProductID), ProductName (ProductName), Employee1ID (Employee1ID), Employee1Name (Employee1Name), Employee2ID (Employee2ID), Employee2Name (Employee2Name)
11. Find **all suppliers** including those who have no products associated with them.
Output: SupplierID (SupplierID), SupplierName (SupplierName),
12. Retrieve a list of **all customers** who have **never placed an order**.
Output: CustomerID (CustomerID), CustomerName (CustomerName)
13. Retrieve **customer pairs who placed orders on the same date**.
Output: Customer1ID (Customer1ID), Customer1Name (Customer1Name), Customer2ID (Customer2ID), Customer2Name (Customer2Name), OrderDate (OrderDate)

14. Find **pairs of employees who have both served the same customer** at any point.

Output: CustomerID (CustomerID), CustomerName
(CustomerName), Employee1ID (Employee1ID), Employee1Name
(Employee1Name), Employee2ID (Employee2ID), Employee2Name
(Employee2Name)

What to Submit:

Submit the following file in Zip on Eduko:

- 2024-CS-X.txt